

Original Analysis Case Study

Part 1 : Graphics Analysis

Part 2 : Feature Reduction (Extraction/Selection)

Part 3 : Filling in Missing Values

Part 1 : Graphics Analysis

In this case study, as part of phase I, we will perform exploratory data analysis by graphing the features in the dataset.

The dataset is composed of 10,000 customer's record at a bank. The dataset has a total of 14 features 13 of which can be considered as independent variables and 1 as the dependent variable. The goal is to build a model that can predict whether a customer is likely to stay or exit the bank. The model will predict the dependent variable 'Exited' using the appropriate set of independent variables

'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumberOfProducts', 'HasCrCard', and 'IsActiveMember'.

We will perform model selection and model validation exercises and use the model to make the desired prediction. The accuracy and precision of the model will be analyzed in the next phases of the study.

```
In [3]: # Load Libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import xlrd
```

```
In [5]: #Step 1: Load data into a dataframe
DataFile = "Data/BankCustomers.xlsx"

data = pd.read_excel(DataFile)
```

```
In [6]: # Step 2: check the dimension of the table
print("The dimension of the table is: ", data.shape)
```

The dimension of the table is: (10000, 14)

```
In [7]: #Step 3: Look at the data
print(data.head(5))
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

In [8]: *#Step 5: what type of variables are in the table*

```
print("Describe Data")
print(data.describe())
```

Describe Data

	RowNumber	CustomerId	CreditScore	Age	Tenur
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.00000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.01280
std	2886.89568	7.193619e+04	96.653299	10.487806	2.89217
min	1.00000	1.556570e+07	350.000000	18.000000	0.00000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.00000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.00000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.00000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.00000

	Balance	NumOfProducts	HasCrCard	IsActiveMember
count	10000.000000	10000.000000	10000.00000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

In [9]: *# Step 6a: Summary of object type data*

```
print("Summarized Data")
print(data.describe(include=['O']))
```

Summarized Data

	Surname	Geography	Gender
count	10000	10000	10000
unique	2932	3	2
top	Smith	France	Male
freq	32	5014	5457

```
In [10]: # Step 6b: Summary of numeric type data
print("Summarized Data")
print(data.describe(include=np.number))
```

Summarized Data					
	RowNumber	CustomerId	CreditScore	Age	Tenur
e \					
count	10000.000000	1.000000e+04	10000.000000	10000.000000	10000.000000
0					
mean	5000.500000	1.569094e+07	650.528800	38.921800	5.012800
0					
std	2886.895680	7.193619e+04	96.653299	10.487806	2.892170
4					
min	1.000000	1.556570e+07	350.000000	18.000000	0.000000
0					
25%	2500.750000	1.562853e+07	584.000000	32.000000	3.000000
0					
50%	5000.500000	1.569074e+07	652.000000	37.000000	5.000000
0					
75%	7500.250000	1.575323e+07	718.000000	44.000000	7.000000
0					
max	10000.000000	1.581569e+07	850.000000	92.000000	10.000000
0					

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	76485.889288	1.530200	0.705500	0.515100	
std	62397.405202	0.581654	0.455840	0.499797	
min	0.000000	1.000000	0.000000	0.000000	
25%	0.000000	1.000000	0.000000	0.000000	
50%	97198.540000	1.000000	1.000000	1.000000	
75%	127644.240000	2.000000	1.000000	1.000000	
max	250898.090000	4.000000	1.000000	1.000000	

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

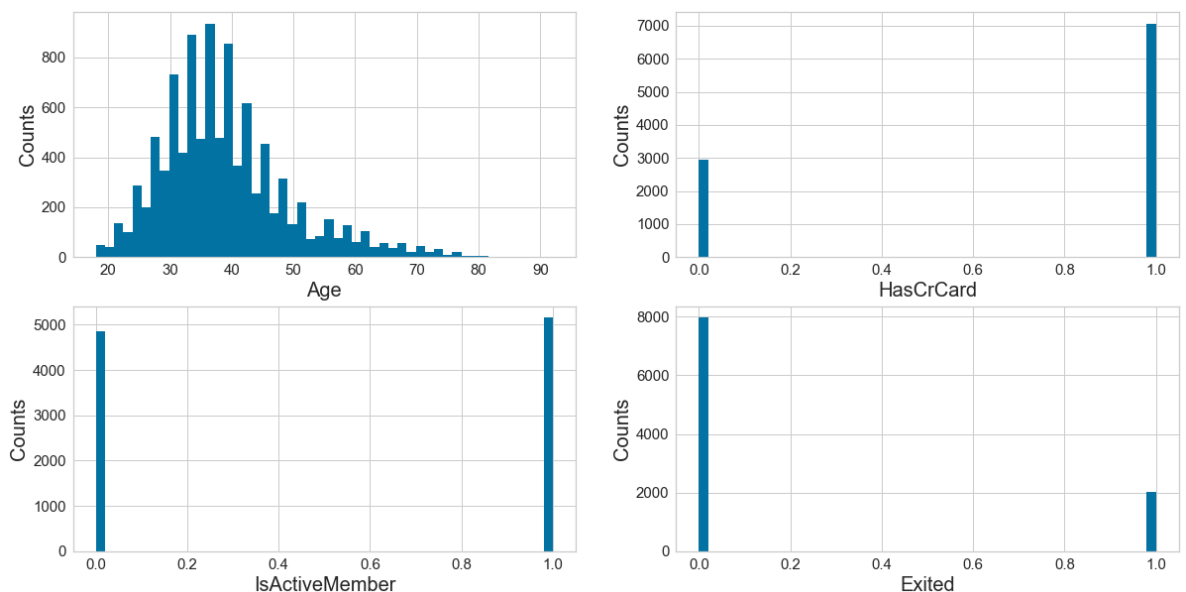
Histogram of ['Age', 'HasCrCard', 'IsActiveMember', 'Exited']

```
In [11]: # set up the figure size
plt.rcParams['figure.figsize'] = (20, 10)

# make subplots
fig, axes = plt.subplots(nrows = 2, ncols = 2)

# Specify the features of interest
num_features = ['Age', 'HasCrCard', 'IsActiveMember', 'Exited']
xaxes = num_features
yaxes = ['Counts', 'Counts', 'Counts', 'Counts']

# draw histograms
axes = axes.ravel()
for idx, ax in enumerate(axes):
    ax.hist(data[num_features[idx]].dropna(), bins=50)
    ax.set_xlabel(xaxes[idx], fontsize=20)
    ax.set_ylabel(yaxes[idx], fontsize=20)
    ax.tick_params(axis='both', labelsize=15)
plt.show()
```



Barchart comparing the number of:

- Exits vs stays
- Males vs. Female
- Has credit card vs does not have credit card
- active members vs inactive members

```

In [12]: # make subplots
fig, axes = plt.subplots(nrows = 2, ncols = 2)





# make the data read to feed into the visulizer
X_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}}).groupby(
    'Exited').size().reset_index(name='Counts')['Exited']
Y_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}}).groupby(
    'Exited').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[0,0].bar(X_Exited, Y_Exited)
axes[0,0].set_title('Exited/Stayed', fontsize=25)
axes[0,0].set_ylabel('Counts', fontsize=20)
axes[0,0].tick_params(axis='both', labelsize=15)

# make the data read to feed into the visulizer
X_Sex = data.groupby('Gender').size().reset_index(name='Counts')['Gender']
Y_Sex = data.groupby('Gender').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[0,1].bar(X_Sex, Y_Sex)
axes[0,1].set_title('Sex', fontsize=25)
axes[0,1].set_ylabel('Counts', fontsize=20)
axes[0,1].tick_params(axis='both', labelsize=15)

X_HasCrCard = data.replace({'HasCrCard': {1: 'Credit Card', 0: 'No Credit Card'}}).groupby('HasCrCard').size().reset_index(name='Counts')['HasCrCard']
Y_HasCrCard = data.replace({'HasCrCard': {1: 'Credit Card', 0: 'No Credit Card'}}).groupby('HasCrCard').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[1,0].bar(X_HasCrCard, Y_HasCrCard)
axes[1,0].set_title('Credit Card', fontsize=25)
axes[1,0].set_ylabel('Counts', fontsize=20)
axes[1,0].tick_params(axis='both', labelsize=15)

X_IsActive = data.replace({'IsActiveMember': {1: 'Active', 0: 'Inactive'}}).groupby('IsActiveMember').size().reset_index(name='Counts')['IsActiveMember']
Y_IsActive = data.replace({'IsActiveMember': {1: 'Active', 0: 'Inactive'}}).groupby('IsActiveMember').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[1,1].bar(X_IsActive, Y_IsActive)
axes[1,1].set_title('Acitivity', fontsize=25)
axes[1,1].set_ylabel('Counts', fontsize=20)
axes[1,1].tick_params(axis='both', labelsize=15)

```



Parallel Coordinate graphe comparing ['Age', 'CreditScore', 'Balance', 'EstimatedSalary']

```

In [13]: # Step 9: Compare variables against those who stayed and those who exited
#set up the figure size
%matplotlib inline
plt.rcParams['figure.figsize'] = (15, 7)
plt.rcParams['font.size'] = 50

# setup the color for yellowbrick visualizer
from yellowbrick.style import set_palette
set_palette('sns_bright')

# import packages
from yellowbrick.features import ParallelCoordinates
# Specify the features of interest and the classes of the target
classes = ['exited', 'stayed']
num_features = ['Age', 'CreditScore', 'Balance', 'EstimatedSalary']

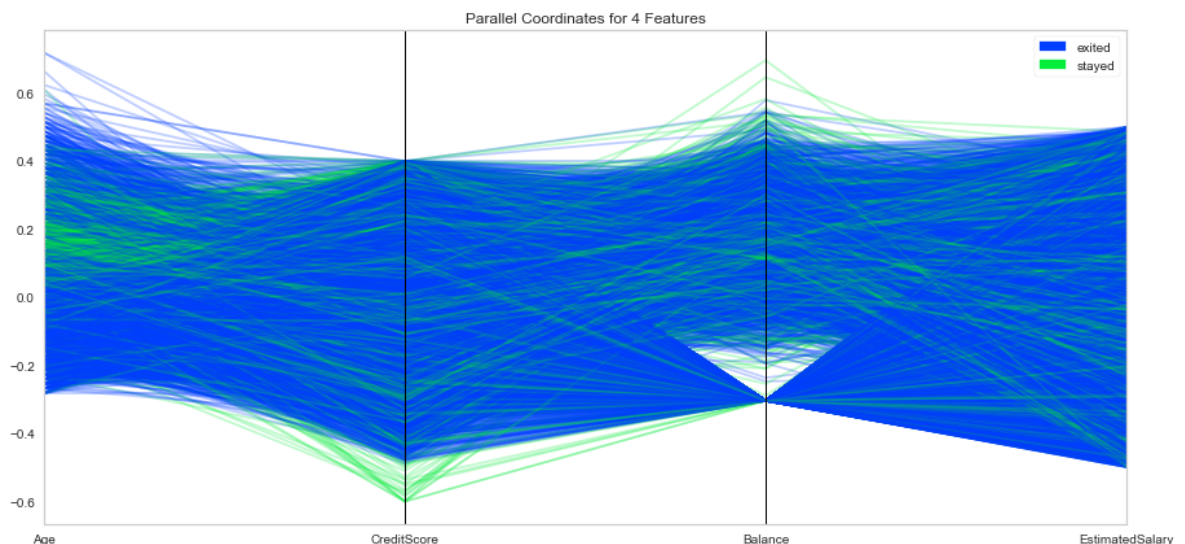
# copy data to a new dataframe
data_norm = data.copy()
# normalize data to 0-1 range
for feature in num_features:
    data_norm[feature] = (data[feature] - data[feature].mean(skipna=True)
    ) / (data[feature].max(skipna=True) - data[feature].min(skipna=True))

# Extract the numpy arrays from the data frame
X = data_norm[num_features].values
y = data.Exited.values

# Instantiate the visualizer
# Instantiate the visualizer
visualizer = ParallelCoordinates(classes=classes, features=num_features)

visualizer.fit(X, y) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data
visualizer.pooof(outpath="images/pcoords2.png") # Draw/show/pooof the data
plt.show();

```



Stacked bar charts showing stays and exits based on:

- Gender
- Has Credit card
- banking activity
- geographic location(Country)

```

In [14]: # Step 10 - stacked bar chart to compare Gender exit/stay numbers
#set up the figure size
%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 20)

# make subplots
fig, axes = plt.subplots(nrows = 2, ncols = 2)

# make the data read to feed into the visulizer
Gender_Stayed = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[data
['Exited']==0]['Gender'].value_counts()
Gender_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[data
['Exited']==1]['Gender'].value_counts()
Gender_Exited = Gender_Exited.reindex(index = Gender_Stayed.index)
# make the bar plot
p1 = axes[0, 0].bar(Gender_Stayed.index, Gender_Stayed.values)
p2 = axes[0, 0].bar(Gender_Exited.index, Gender_Exited.values, bottom=Gen
der_Stayed.values)
axes[0, 0].set_title('Gender Stayed/Exited', fontsize=25)
axes[0, 0].set_ylabel('Counts', fontsize=20)
axes[0, 0].tick_params(axis='both', labelsize=15)
axes[0, 0].legend((p1[0], p2[0]), ('Stayed', 'Exited'), fontsize = 15)

# make the data read to feed into the visulizer
HasCrCard_Stayed = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[d
ata['Exited']==0]
HasCrCard_Stayed = HasCrCard_Stayed.replace({'HasCrCard': {1: 'CreditCar
d', 0: 'No CreditCard'}})['HasCrCard'].value_counts()

HasCrCard_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[d
ata['Exited']==1]
HasCrCard_Exited = HasCrCard_Exited.replace({'HasCrCard': {1: 'CreditCar
d', 0: 'No CreditCard'}})['HasCrCard'].value_counts()
HasCrCard_Exited = HasCrCard_Exited.reindex(index = HasCrCard_Stayed.inde
x)
# make the bar plot
p3 = axes[0, 1].bar(HasCrCard_Stayed.index, HasCrCard_Stayed.values)
p4 = axes[0, 1].bar(HasCrCard_Exited.index, HasCrCard_Exited.values, bott
om=HasCrCard_Stayed.values)
axes[0, 1].set_title('HasCrCard Stayed/Exited', fontsize=25)
axes[0, 1].set_ylabel('Counts', fontsize=20)
axes[0, 1].tick_params(axis='both', labelsize=15)
axes[0, 1].legend((p3[0], p4[0]), ('Stayed', 'Exited'), fontsize = 15)

# make the data read to feed into the visulizer
IsActive_Stayed = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[da
ta['Exited']==0]
IsActive_Stayed = IsActive_Stayed.replace({'IsActiveMember': {1: 'Active'
, 0: 'Inactive'}})['IsActiveMember'].value_counts()

IsActive_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[da
ta['Exited']==1]
IsActive_Exited = IsActive_Exited.replace({'IsActiveMember': {1: 'Active'
, 0: 'Inactive'}})['IsActiveMember'].value_counts()
IsActive_Exited = IsActive_Exited.reindex(index = IsActive_Stayed.index)
# make the bar plot

```

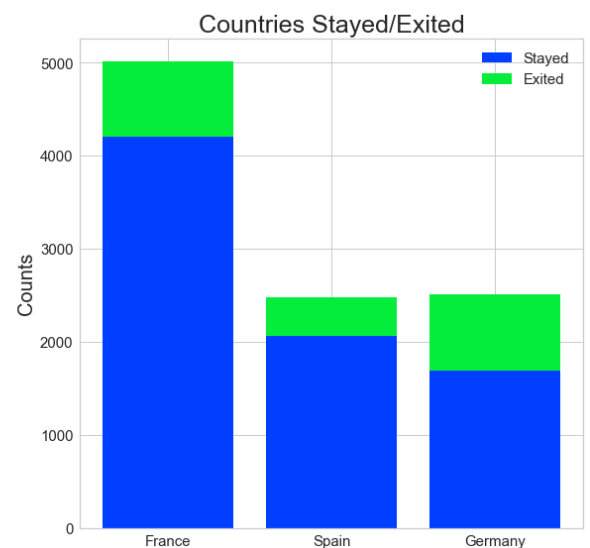
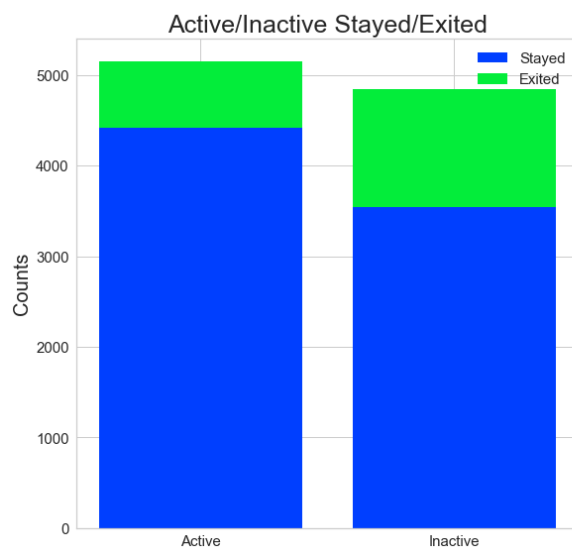
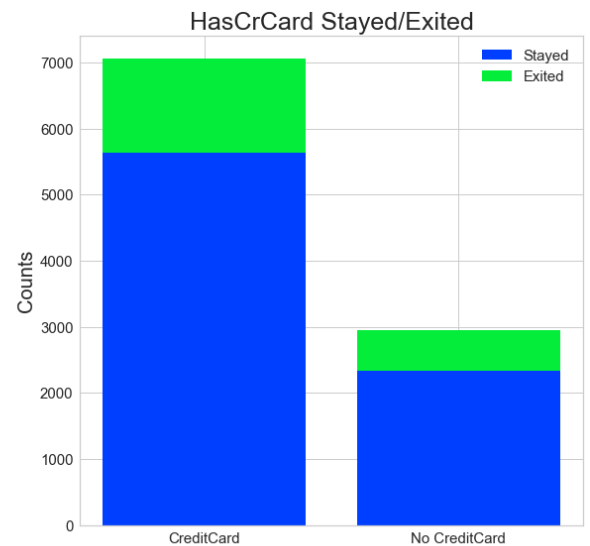
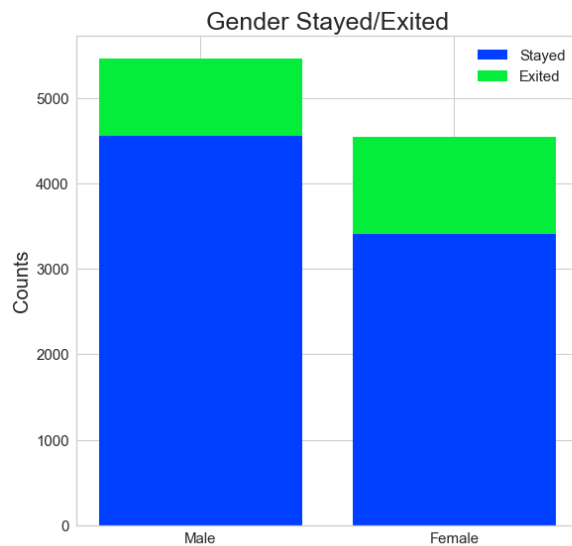
```

p4 = axes[1,0].bar(IsActive_Stayed.index, IsActive_Stayed.values)
p5 = axes[1,0].bar(IsActive_Exited.index, IsActive_Exited.values, bottom=
IsActive_Stayed.values)
axes[1,0].set_title('Active/Inactive Stayed/Exited', fontsize=25)
axes[1,0].set_ylabel('Counts', fontsize=20)
axes[1,0].tick_params(axis='both', labelsize=15)
axes[1,0].legend((p4[0], p5[0]), ('Stayed', 'Exited'), fontsize = 15)

# make the data read to feed into the visulizer
Country_Stayed = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[data[
'Exited']==0]['Geography'].value_counts()

Country_Exited = data.replace({'Exited': {1: 'Exited', 0: 'Stayed'}})[data[
'Exited']==1]['Geography'].value_counts()
Country_Exited = Country_Exited.reindex(index = Country_Stayed.index)
# make the bar plot
p6 = axes[1,1].bar(Country_Stayed.index, Country_Stayed.values)
p7 = axes[1,1].bar(Country_Exited.index, Country_Exited.values, bottom=Co
untry_Stayed.values)
axes[1,1].set_title('Countries Stayed/Exited', fontsize=25)
axes[1,1].set_ylabel('Counts', fontsize=20)
axes[1,1].tick_params(axis='both', labelsize=15)
axes[1,1].legend((p6[0], p7[0]),('Stayed', 'Exited'), fontsize = 15)
plt.show()

```



Part 2 : Feature Reduction (Extraction/Selection)

```
In [15]: #Step 1: Load data into a dataframe
DataFile = "Data/BankCustomers.xlsx"

data = pd.read_excel(DataFile)
```

```
In [16]: data.columns
```

```
Out[16]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
                'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited'],
                dtype='object')
```

```
In [17]: # Step 11- remove unrelated columns
data = data.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
```

```
In [18]: data.columns
```

```
Out[18]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',  
              'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',  
              'Exited'],  
              dtype='object')
```

```
In [19]: # Step 12 - Onehot code Geography  
from sklearn.preprocessing import LabelBinarizer, MultiLabelBinarizer  
  
feature = np.array(data['Geography'])  
one_hot = LabelBinarizer()  
  
one_hot.fit_transform(feature)
```

```
Out[19]: array([[1, 0, 0],  
               [0, 0, 1],  
               [1, 0, 0],  
               ...,  
               [1, 0, 0],  
               [0, 1, 0],  
               [1, 0, 0]], dtype=int32)
```

```
In [20]: one_hot.classes_
```

```
Out[20]: array(['France', 'Germany', 'Spain'], dtype='<U7')
```

```
In [21]: dummies = pd.get_dummies(feature)  
dummies.head()
```

```
Out[21]:
```

	France	Germany	Spain
0	1	0	0
1	0	0	1
2	1	0	0
3	1	0	0
4	0	0	1

```
In [22]: # Drop Geography column  
data = data.drop(['Geography'],axis=1)
```

```
In [23]: # Add dummies
data[dummies.columns] = dummies
data.head()
```

```
Out[23]:
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActi
0	619	Female	42	2	0.00	1	1	
1	608	Female	41	1	83807.86	1	0	
2	502	Female	42	8	159660.80	3	1	
3	699	Female	39	1	0.00	2	0	
4	850	Female	43	2	125510.82	1	1	

```
In [24]: # one-hot code Gender
feature = np.array(data['Gender'])
one_hot = LabelBinarizer()

one_hot.fit_transform(feature)
dummies = pd.get_dummies(feature)
dummies
```

```
Out[24]:
```

	Female	Male
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0
...
9995	0	1
9996	0	1
9997	1	0
9998	0	1
9999	1	0

10000 rows × 2 columns

```
In [25]: #drop Gender and add dummies
data = data.drop(['Gender'],axis=1)
data[dummies.columns] = dummies
data.head()
```

```
Out[25]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	

```
In [26]: # Drop spain and male to avoid dummy trap
data = data.drop(['Male', 'Spain'],axis=1)
data.head()
```

```
Out[26]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	

```
In [64]: # Move the dependent variable column to the last position.

Exited = data.replace({'Exited': {1: 'Existed', 0: 'Stayed'}})['Exited']
```

```
In [97]: data = data.drop(['Exited'],axis=1)
data['Exited'] = Exited
data.head()
```

```
Out[97]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Exited
0	619	42	2	0.00	1	1		Stayed
1	608	41	1	83807.86	1	0		Stayed
2	502	42	8	159660.80	3	1		Existed
3	699	39	1	0.00	2	0		Stayed
4	850	43	2	125510.82	1	1		Existed

```
In [99]: # Step 13 - Set up independent variable and dependent variables and perform feature reduction
Independents = data.iloc[:, :-1].values
print(type(Independents))
Dependent = data.iloc[:, -1].values
print(Dependent)
X = Independents
y = Dependent
```

```
<class 'numpy.ndarray'>
['Existed' 'Stayed' 'Existed' ... 'Existed' 'Existed' 'Stayed']
```

```
In [100]: data.shape
```

```
Out[100]: (10000, 12)
```

```
In [101]: X.shape
```

```
Out[101]: (10000, 11)
```

```
In [102]: y.shape
```

```
Out[102]: (10000,)
```

```
In [103]: # Attempt at feature reduction using PCA Before feature scaling
#Load libraries
from sklearn.preprocessing import StandardScaler

from sklearn.decomposition import PCA

# Create a PCA that will retain 99% of variance
pca = PCA(n_components=0.99, whiten=True)

# Conduct PCA
features_pca = pca.fit_transform(X)

# Show results
print("Original number of features:", X.shape[1])
print("Reduced number of features:", features_pca.shape[1])
```

```
Original number of features: 11
Reduced number of features: 2
```



```
In [104]: # Feature scaling will normalize all variable to the same scale
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X = sc.fit_transform(X)
print(X)

[[-0.32622142  0.29351742 -1.04175968 ...  0.99720391 -0.57873591
  1.09598752]
 [-0.44003595  0.19816383 -1.38753759 ... -1.00280393 -0.57873591
  1.09598752]
 [-1.53679418  0.29351742  1.03290776 ...  0.99720391 -0.57873591
  1.09598752]
 ...
 [ 0.60498839 -0.27860412  0.68712986 ...  0.99720391 -0.57873591
  1.09598752]
 [ 1.25683526  0.29351742 -0.69598177 ... -1.00280393  1.72790383
 -0.91241915]
 [ 1.46377078 -1.04143285 -0.35020386 ...  0.99720391 -0.57873591
  1.09598752]]
```

```
In [105]: # Attempt at feature reduction using PCA After feature scaling
#Load libraries
from sklearn.preprocessing import StandardScaler

from sklearn.decomposition import PCA

# Create a PCA that will retain 99% of variance
pca = PCA(n_components=0.99, whiten=True)

# Conduct PCA
features_pca = pca.fit_transform(X)

# Show results
print("Original number of features:", X.shape[1])
print("Reduced number of features:", features_pca.shape[1])

Original number of features: 11
Reduced number of features: 11
```

Part 3 : Filling in Missing Values

Summary of parts 1 and 2: We have performed feature reduction and scaled the independent variables. The X and y variables are the independent variables dataset and the dependent variables respectively. The value of 0 or 1 for the depended variable has been converted to 'Stayed' and 'Exited' respectively in anticipation of using logistic regression classifier for modeling.

- Split_Train_Test

- Model Selection and Evaluation

```
In [107]: import pandas as pd
import yellowbrick

import warnings
warnings.filterwarnings("ignore")
```

```
In [111]: print("Indpenden variables matrix:\n")
print(X)
```

Indpenden variables matrix:

```
[[-0.32622142  0.29351742 -1.04175968 ...  0.99720391 -0.57873591
  1.09598752]
 [-0.44003595  0.19816383 -1.38753759 ... -1.00280393 -0.57873591
  1.09598752]
 [-1.53679418  0.29351742  1.03290776 ...  0.99720391 -0.57873591
  1.09598752]
 ...
 [ 0.60498839 -0.27860412  0.68712986 ...  0.99720391 -0.57873591
  1.09598752]
 [ 1.25683526  0.29351742 -0.69598177 ... -1.00280393  1.72790383
 -0.91241915]
 [ 1.46377078 -1.04143285 -0.35020386 ...  0.99720391 -0.57873591
  1.09598752]]
```

```
In [112]: print("Dependent variable array:\n")
print(y)
```

Dependent variable array:

```
['Existed' 'Stayed' 'Existed' ... 'Existed' 'Existed' 'Stayed']
```

Step 14 - Split the dataset to 30% test set and 70% training dataset

```
In [114]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =0.3,
random_state=0)

# number of samples in each set
print("Total sample in dataset: ", X.shape[0])

print("No. of samples in training set: ", X_train.shape[0])
print("No. of samples in validation set:", X_test.shape[0])
```

Total sample in dataset: 10000
No. of samples in training set: 7000
No. of samples in validation set: 3000

```
In [115]: print(y_train.shape)
```

```
(7000,)
```

```
In [116]: print(y_val.shape)
```

```
(3000,)
```

```
In [117]: # stayed and exited  
print('\n')  
print('No. of customer who stayed and exited in the training set:')  
print(pd.Series(y_train).value_counts())
```

```
No. of customer who stayed and exited in the training set:
```

```
Stayed      5584
```

```
Existed     1416
```

```
dtype: int64
```

```
In [118]: print('\n')  
print('No. of customer who stayed and exited in the validation set:')  
print(pd.Series(y_val).value_counts())
```

```
No. of customer who stayed and exited in the validation set:
```

```
Stayed      2379
```

```
Existed      621
```

```
dtype: int64
```

Step 15 - Model evaluation and metrics

Create a logistics regression model

```
In [119]: from sklearn.linear_model import LogisticRegression  
  
from yellowbrick.classifier import ConfusionMatrix  
from yellowbrick.classifier import ClassificationReport  
from yellowbrick.classifier import ROCAUC  
  
# Instantiate the classification model  
model = LogisticRegression()
```

Define class for 'Exited' and 'stayed' to create confusion metrix and fit it into the trainign sets. Then display the confusion metric


```
In [120]: #The ConfusionMatrix visualizer takes a model
classes = ['Exited', 'Stayed']
cm = ConfusionMatrix(model, classes=classes, percent=False)

#Fit fits the passed model. This is unnecessary if you pass the visualizer a pre-fitted model
cm.fit(X_train, y_train)

#To create the ConfusionMatrix, we need some test data. Score runs predict() on the data
#and then creates the confusion_matrix from scikit learn.
cm.score(X_val, y_val)

# change fontsize of the labels in the figure
for label in cm.ax.texts:
    label.set_size(20)

#How did we do?
cm.poof()
```



```
Out[120]: <matplotlib.axes._subplots.AxesSubplot at 0x10be7bb0>
```

Precision, Recall, and F1 Score metrics:

```
In [95]: %matplotlib inline
plt.rcParams['figure.figsize'] = (15, 7)
plt.rcParams['font.size'] = 20

# Instantiate the visualizer
visualizer = ClassificationReport(model, classes=classes)

visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_val, y_val) # Evaluate the model on the test data
g = visualizer.poof()

# ROC and AUC
#Instantiate the visualizer
visualizer = ROCAUC(model)

visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_val, y_val) # Evaluate the model on the test data
g = visualizer.poof()
```

